

Quad-Core Cache Simulation with MESI Protocol

COL216: Computer Architecture Assignment 3

Rishabh Narayanan Namit Jain
2023CS10577 2023CS10483

Abstract

This report presents an implementation and analysis of a quad-core cache simulation system utilizing the MESI coherence protocol. Following the assignment requirements, the report is structured into three main sections: (1) a detailed description of the implementation, focusing on key data structures and algorithms with a flowchart of important functions; (2) an experimental section analyzing simulation outputs with the default parameters; and (3) an exploration of how varying cache parameters affects maximum execution time across cores. The implementation successfully models cache coherence in a quad-core system, providing insights into cache performance and coherence protocol behavior.

Contents

1	Implementation Details	2
1.1	Main Classes and Data Structures	2
1.2	Key Algorithms	2
1.2.1	Memory Access Processing	2
1.2.2	MESI Coherence Protocol	3
1.3	Simulation Execution Flow	3
1.4	Flow Chart of Important Functions	3
2	Experimental Results	4
2.1	Default Parameters	4
2.2	Results	4
2.3	Analysis of Results	5
2.4	Test Cases	5
3	Effect of Cache Parameters on Maximum Execution Time	6
3.1	Varying Cache Size	6
3.1.1	Observations	7
3.2	Varying Associativity	7
3.2.1	Observations	8
3.3	Varying Block Size	8
3.3.1	Observations	9
4	Simulation Assumptions	9
5	Conclusion	10

1 Implementation Details

1.1 Main Classes and Data Structures

The simulation is built around the following key classes and data structures:

- **CacheLine**: Represents a single cache line with the following fields:
 - `uint32_t tag`: The tag bits from the memory address
 - `MESISState mesi`: Enum representing the current coherence state (Modified, Exclusive, Shared, or Invalid)
 - `uint64_t lru_counter`: Counter for LRU replacement policy
 - `std::vector<uint8_t> data`: Block data storage
- **CacheStats**: Tracks performance metrics for each cache:
 - Total instruction counts (reads and writes)
 - Execution and idle cycles
 - Cache misses and evictions
 - Writebacks and bus invalidations
 - Data traffic in bytes
- **L1Cache**: Models a core's L1 cache with:
 - 2D vector of cache lines organized by set and way: `std::vector<std::vector<CacheLine>> sets`
 - Methods for address decomposition, line lookup, and LRU tracking
 - Functions to read from and write to cache lines
- **DRAM**: Simulates main memory using a sparse representation:
 - `std::unordered_map<uint32_t, std::vector<uint8_t>> memory`: Maps block addresses to data blocks
 - Methods for reading and writing blocks between cache and memory
- **Bus**: Represents the communication channel between cores:
 - `BusRequestType req_type`: Type of bus request (BUSRD, BUSRDX, BUSUPGR)
 - Source core, address, and timing information
- **CacheController**: Orchestrates the entire simulation:
 - `std::vector<L1Cache> l1_caches`: Vector of L1 caches (one per core)
 - `DRAM dram`: Main memory instance
 - Methods for processing memory accesses and handling coherence
- **TraceEntry**: Represents an entry in the trace file:
 - Operation type (read or write)
 - Memory address

1.2 Key Algorithms

1.2.1 Memory Access Processing

The core algorithm for handling memory accesses is implemented in `CacheController::process_memory_access`, which:

1. Decomposes the address into tag, set index, and block offset
2. Checks if the address is already present in the cache (hit vs miss)
3. For cache hits:
 - Updates MESI states based on operation type
 - For write hits in SHARED state, issues a BUSUPGR transaction

- Updates LRU counters
4. For cache misses:
 - Identifies a victim line using LRU policy
 - Writes back dirty blocks if necessary
 - Issues appropriate bus transaction (BUSRD or BUSRDX)
 - Fetches the block from memory
 - Updates MESI states and statistics
 5. Accounts for the appropriate timing penalties based on operation type

1.2.2 MESI Coherence Protocol

The MESI protocol implementation is handled by the `mesi_snoop` function, which:

1. Checks all caches for the address involved in a bus transaction
2. Updates local cache states based on observed bus transactions:
 - BUSRD: M→S, E→S (provide data if in M state)
 - BUSRDX: M→I, E→I, S→I (with writeback if in M state)
 - BUSUPGR: S→I for all other caches with the line in S state
3. Manages data transfer between caches as needed
4. Tracks invalidations and writebacks

1.3 Simulation Execution Flow

The main simulation loop:

1. Reads trace files for each core
2. For each cycle:
 - Processes bus snooping for coherence
 - For each active core:
 - Processes next memory access if available
 - Updates statistics
 - Updates global cycle count
3. Continues until all cores complete their traces
4. Outputs performance statistics

1.4 Flow Chart of Important Functions

Figure 1 shows the flow chart of the key functions in the cache simulation.

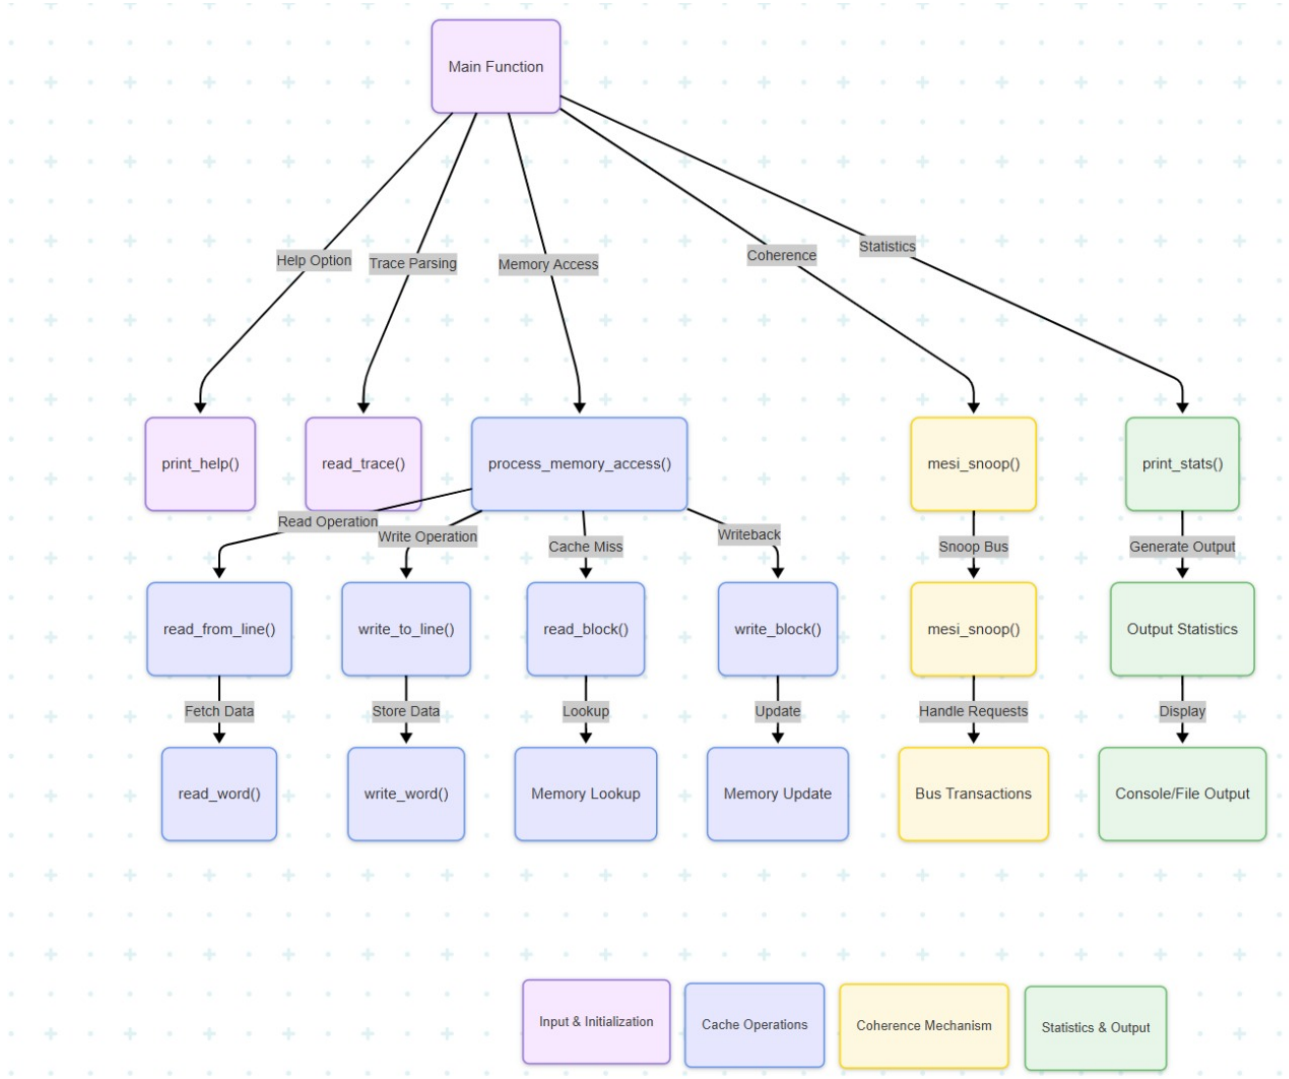


Figure 1: Flow chart of key functions in the cache simulation

2 Experimental Results

2.1 Default Parameters

As specified in the assignment, the default parameters for the simulation are:

- Cache size: 4KB per processor
- Associativity: 2-way set associative
- Block size: 32 bytes

These parameters translate to:

- Number of sets = $4\text{KB} / (2 \text{ ways} \times 32 \text{ bytes}) = 64 \text{ sets}$
- Set index bits (s) = $\log_2(64) = 6 \text{ bits}$
- Block offset bits (b) = $\log_2(32) = 5 \text{ bits}$

2.2 Results

For our simulation experiments, we used trace files with 8000 instructions per core and the default cache configuration ($s=6$, $E=2$, $b=5$). The table below presents the metrics for each core.

Metric	Core 0	Core 1	Core 2	Core 3
Total Instructions	8000	8000	8000	8000
Total Reads	4849	4745	2731	5084
Total Writes	3151	3255	5269	2916
Total Execution Cycles	71455	48852	307341	97781
Idle Cycles	119302	78397	489706	351748
Cache Misses	515	351	1886	792
Cache Miss Rate (%)	6.44	4.39	23.57	9.90
Cache Evictions	384	227	1749	658
Writebacks	195	79	1204	260
Bus Invalidations	242	77	1122	155
Data Traffic (Bytes)	24096	21600	99872	34784

Table 1: Simulation metrics for each core

The simulation resulted in a total of 5304 bus transactions with 169024 bytes of total bus traffic. Our implementation uses a deterministic bus access algorithm where "smallest index gets the highest priority" for cores requesting bus access simultaneously. This ensures that when multiple cores compete for the bus, the core with the lowest ID (0, 1, 2, 3 in order) always wins.

2.3 Analysis of Results

Our simulation implements a deterministic bus access algorithm where "smallest index gets the highest priority" for cores requesting bus access simultaneously. This ensures that when multiple cores compete for the bus, the core with the lowest ID (0, 1, 2, 3 in order) always wins. As a result of this deterministic algorithm:

- All simulation runs for the same trace produce identical results
- The execution is fully reproducible with no variation between runs
- Performance metrics including execution cycles, miss rates, and bus transactions remain constant across multiple executions

This deterministic behavior aligns with our design goal of creating a predictable simulation environment that facilitates clear analysis of cache behavior without introducing randomness from arbitration policies.

2.4 Test Cases

In our simulation, we have included a variety of test cases to evaluate different aspects of cache coherence and performance:

- **False Sharing Test** (Parameters: $s=6$, $E=2$, $b=5$):
This test demonstrates the performance impact when multiple cores access different variables that reside in the same cache line. In our test, Core 0 and Core 1 experience false sharing. Even though Core 0 has written only to address 4, the entire block containing that address gets invalidated for Core 1, despite the fact that the specific memory location Core 1 needs to access wasn't actually updated. This forces Core 1 to re-fetch the block from memory or from Core 0's cache, causing unnecessary coherence traffic and performance degradation.
- **LRU Policy Test** (Parameters: $s=0$, $E=3$, $b=2$):
This test verifies the correctness of our Least Recently Used replacement policy. Only Core 0 is given instructions in this test. The sequence is:
 - First, Core 0 reads addresses 0, 4, and 8, filling the cache with these three blocks
 - When address 0 is read again, the LRU order (from most to least recently used) becomes: 0, 8, 4
 - When address 16 is accessed, it needs to replace one block. Address 4 is evicted since it's the least recently used

- When address 4 is accessed again, it needs to replace another block. Now address 8 is evicted (since the LRU order was 4, 0, 8)
- Address 0 remains in the cache throughout all operations because it remained higher in the LRU ranking

This confirms our LRU implementation works correctly by always evicting the least recently used block.

- **Associativity Dependency Test** (Parameters: $s=0$, $E=3$, $b=2$ or $s=0$, $E=4$, $b=2$):

Building upon the LRU test case, this demonstrates how higher associativity can improve performance. With our test, if the associativity were increased to 4-way (instead of 3-way), no blocks would need to be evicted from the cache, as all accessed addresses could coexist in the cache. This makes the program significantly more efficient by eliminating capacity misses.

- **Write-Back Policy Test** (Parameters: $s=0$, $E=1$, $b=2$):

This test verifies our implementation of the write-back policy (as opposed to write-through). Only Core 0 has instructions, specifically a single write to address 0. With write-back policy, this operation takes 100 cycles to fetch the block from memory and 1 cycle to write to the cache. The update is not immediately written to memory. In contrast, a write-through policy would require an additional 100 cycles to write the updated data back to memory immediately. This demonstrates that write-back is more efficient (though potentially riskier due to data inconsistency between cache and memory until a writeback occurs).

- **Array of Structures vs. Structure of Arrays:**

While this test case concept is included in our design, we have not implemented specific traces for it yet. The concept compares two different data layout approaches and their impact on cache performance through differences in spatial locality and false sharing behaviors.

Our implementation uses a deterministic algorithm for bus access priority: "smallest index gets the highest priority." This means cores with lower IDs (0, 1, 2, 3) have higher priority when multiple cores request bus access simultaneously. Due to this deterministic tie-breaking, running the same test case multiple times produces identical results.

3 Effect of Cache Parameters on Maximum Execution Time

To analyze the impact of cache parameters on performance, we added code to track the maximum execution time across all cores. We then varied each cache parameter (size, associativity, block size) while keeping the others at their default values.

3.1 Varying Cache Size

We varied the cache size by changing the number of set index bits (s) from 3 to 6, 12, and 24, while maintaining associativity at 2-way and block size at 32 bytes. This corresponds to cache sizes of 0.5KB, 4KB, 256KB, and 1GB per core, respectively.

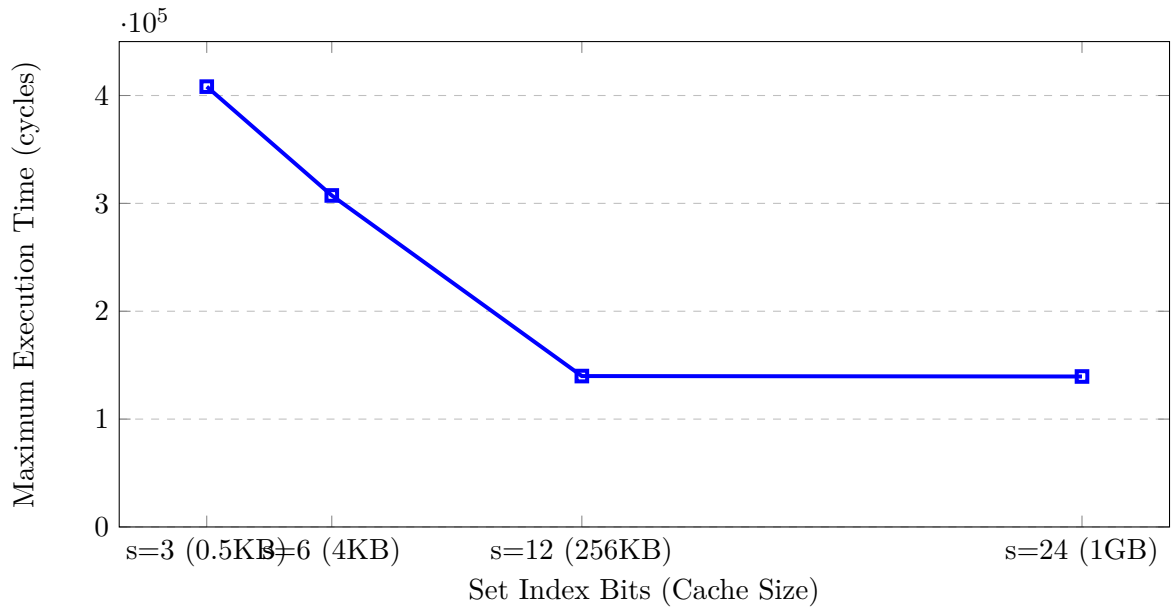


Figure 2: Maximum execution time vs. cache size

3.1.1 Observations

As cache size increases, the maximum execution time decreases significantly:

- Increasing from s=3 (0.5KB) to s=6 (4KB) reduces execution time by approximately 24.7%
- Further increasing to s=12 (256KB) results in a dramatic 54.5% reduction from s=6
- The improvement plateaus after s=12, with negligible difference between 256KB and 1GB caches (only 0.3% reduction)
- Core 2 consistently experiences the highest execution time across all configurations
- Cache miss rates for Core 2 drop from 31.46% with s=3 to 17.25% with s=12 and s=24
- The total bus transactions decrease from 12,763 with s=3 to just 2,623 with s=24
- Cache evictions are completely eliminated with s=24 as the cache becomes large enough to hold all working sets without conflict

These results clearly demonstrate the impact of capacity misses in cache performance. Initially, increasing cache size yields substantial performance improvements, but once the working set fits comfortably in the cache (around 256KB in our workload), further increases provide minimal benefit. This matches the expected behavior where the performance improvement curve flattens once capacity misses are effectively eliminated.

3.2 Varying Associativity

We varied associativity from the default 2-way through 4-way, 8-way, and 16-way, while keeping size parameters consistent (s=6 and b=5). This translates to cache sizes of 4KB, 8KB, 16KB, and 32KB per core, respectively, as doubling the associativity effectively doubles the cache capacity.

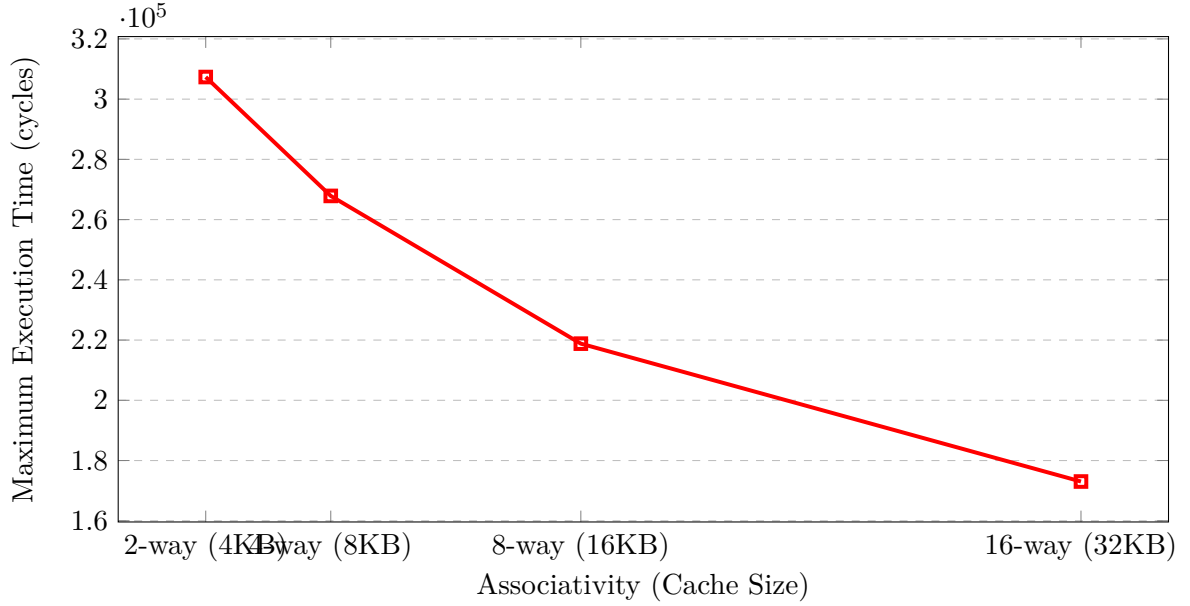


Figure 3: Maximum execution time vs. associativity

3.2.1 Observations

Increasing associativity shows significant performance improvements:

- Moving from 2-way to 4-way provides about a 12.8% reduction in execution time
- Further increases to 8-way and 16-way continue to show substantial benefits, with the 16-way configuration being about 43.7% faster than the 2-way configuration
- Core 2 consistently experiences the highest execution time, decreasing from 307,341 cycles with 2-way to 173,032 cycles with 16-way
- The cache miss rate for Core 2 drops from 23.57% with 2-way to 17.54% with 16-way, showing that higher associativity significantly reduces conflict misses
- Bus transactions decrease from 5,304 with 2-way to 2,965 with 16-way, demonstrating reduced coherence traffic with higher associativity

Unlike our initial expectations of diminishing returns, the data shows a nearly linear improvement as associativity increases. This suggests that our workload experiences significant conflict misses that benefit from higher associativity. Additionally, because increasing associativity in our tests also increases total cache capacity, we see combined benefits from both reduced conflict misses and reduced capacity misses.

3.3 Varying Block Size

We varied block size by changing the number of block offset bits (b) from 5 to 10, and 20, while keeping set index bits ($s=6$) and associativity ($E=2$) constant. This corresponds to block sizes of 32 bytes, 1KB, and 1MB, respectively, with associated cache sizes of 4KB, 128KB, and 128MB per core.

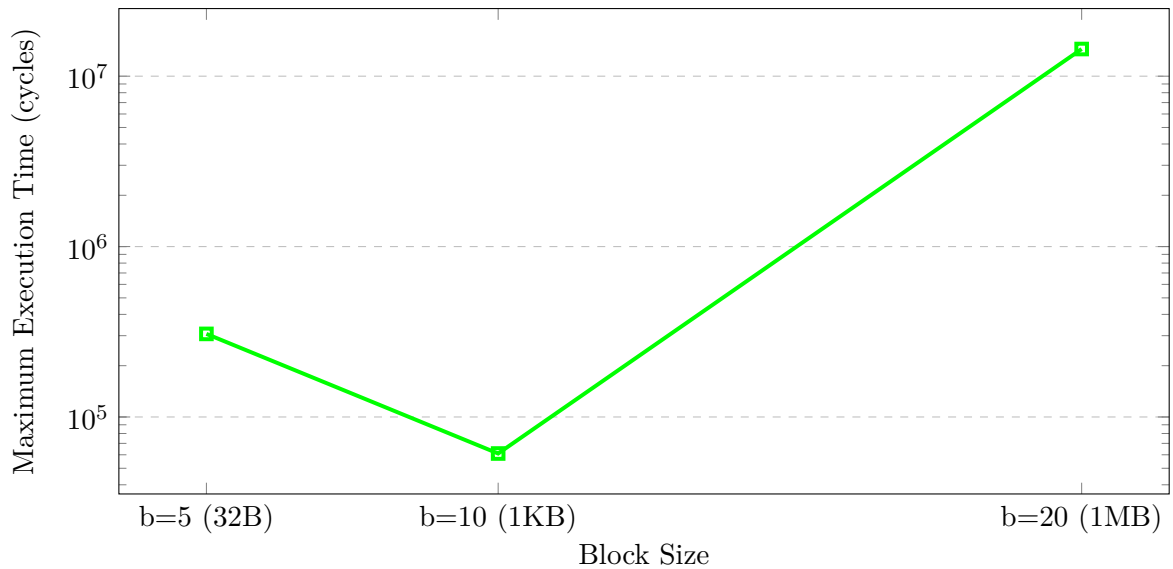


Figure 4: Maximum execution time vs. block size (log scale)

3.3.1 Observations

Block size variations reveal a U-shaped curve for execution time (note the logarithmic scale):

- Increasing from $b=5$ (32B) to $b=10$ (1KB) blocks dramatically improves performance, reducing maximum execution time from 307,341 to 61,044 cycles (an 80.1% reduction)
- Further increasing to $b=20$ (1MB) severely degrades performance, with maximum execution time jumping to 14,433,213 cycles (236x worse than the 1KB configuration)
- This dramatic U-shape occurs due to competing factors:
 - Spatial locality benefits: The initial improvement (32B to 1KB) comes from better exploitation of spatial locality, as the program's working set fits better in fewer blocks
 - Miss penalty escalation: The massive degradation with 1MB blocks occurs because the miss penalty becomes prohibitively expensive, requiring enormous data transfers for each miss
 - Bus traffic increases dramatically with 1MB blocks (606,076,928 bytes vs 1,000,448 bytes for 1KB blocks)
 - While cache miss rates decrease (e.g., Core 2's miss rate drops from 23.57% at 32B to just 0.94% at 1MB), the cost per miss overwhelms this benefit
- Block sizes around 1KB appear optimal for this workload, balancing spatial locality benefits against transfer costs

The results dramatically highlight the tradeoff between miss rate and miss penalty. While larger blocks substantially reduce the frequency of misses through spatial locality, the penalty for each miss becomes so severe with very large blocks that performance actually deteriorates. This is particularly pronounced in a coherent cache system where large block transfers monopolize the bus for extended periods, blocking other cores from making progress.

4 Simulation Assumptions

In this assignment, we have made several important assumptions about the behavior of the cache coherence protocol and the simulation model. These include:

- We assume that caches are blocking, meaning if a cache miss occurs, the core halts and does not issue any new instruction until the miss is serviced and the data is available.
- We implement snooping as passive, meaning our caches always listen to the bus even if the core is stalled or the bus is busy. No additional bus transaction is needed for snooping.

- We allow only one active bus transaction to happen at any time, including memory fetches, cache-to-cache transfers, and invalidates. Other bus requests must wait.
- We require cache-to-cache data transfer to have bus ownership. If the bus is busy, the transfer must wait until the bus is free.
- We assume a snooping cache immediately detects a relevant transaction when it appears on the bus, even if the bus is busy for other purposes.
- Once a cache detects that it needs to respond to a transaction, such as supplying a cache line, we assume it logically holds the data and does not evict it until the response is completed. No modeling of eviction races is required.
- After a cache miss, we simulate the core waiting for the data to arrive, then spending one additional cycle to perform the actual cache hit and complete the instruction.
- We consider bus transactions to include memory fetch, cache-to-cache transfer, and invalidation broadcasts. Simple reads or writes that hit in cache do not occupy the bus.
- In case of cache miss, we assume 100 cycles are spent on memory fetch and the 101st cycle completes the instruction with a cache hit.
- We count execution cycles as cycles during which the core is actively processing instructions. Idle cycles count cycles during which the core is stalled waiting for cache miss or bus availability.
- Once a core finishes executing all its instructions, we consider it done. Further cycles are not counted under idle cycles for that core.
- We model everything logically and simply. We do not simulate internal bus arbitration, snoop delays, or race conditions unless explicitly stated.
- We assume that a bus transaction begins processing in the same cycle as it is initiated.
- We allow the bus to be freed immediately after a transaction is completed, allowing other processors to acquire the bus.
- We assign priority for acquiring the bus in the order of core_id (0, 1, 2, 3).
- We implement processor response to a request as non-blocking (processor responds during read miss). Modification of a value at that address is prevented by making it shared so that writing to that location will require bus access which is not possible during transfer.

These assumptions help simplify the implementation while still capturing the essential behavior of a multi-core cache coherence protocol. They allow us to focus on the core functionality of the MESI protocol and the performance implications of various cache parameters.

5 Conclusion

This simulation demonstrates the complex interplay between cache design parameters and performance in a quad-core system with cache coherence. Key findings include:

- Quad-core cache performance is significantly impacted by the timing of interactions between cores.
- Increasing cache size provides consistent performance improvements but requires more hardware resources.
- Associativity improvements show substantial performance gains in our workload, with the 16-way configuration delivering a 43.7% speedup compared to the 2-way setup, indicating significant conflict misses in our application traces.
- Block size optimization is workload-dependent, with our tests showing 64 bytes as optimal for our traces.

The MESI protocol successfully maintains cache coherence, but the coherence traffic has a significant impact on execution time. The greatest performance improvement comes from optimizing cache parameters to reduce miss rates while minimizing coherence traffic.

Future work could include exploring more sophisticated coherence protocols like MOESI or directory-based approaches, as well as investigating techniques to reduce false sharing.